



# COMP9311: Database Systems

**Term 3 2022**

**Week 4 (SQL)**

**By Helen Paik, CSE UNSW**

**Textbook: Chapters 6 and 7**

**Disclaimer: the course materials are sourced from**

- previous offerings of COMP9311 and COMP3311
- Prof. Werner Nutt on Introduction to Database Systems (<http://www.inf.unibz.it/~nutt/Teaching/IDBs1011/>)

## MotherChild

<b>mother</b>	<b>child</b>
Lisa	Mary
Lisa	Greg
Anne	Kim
Anne	Phil
Mary	Andy
Mary	Rob

## FatherChild

<b>father</b>	<b>child</b>
Steve	Frank
Greg	Kim
Greg	Phil
Frank	Andy
Frank	Rob

## Person

<b>name</b>	<b>age</b>	<b>income</b>
Andy	27	21
Rob	25	15
Mary	55	42
Anne	50	35
Phil	26	30
Greg	50	40
Frank	60	20
Kim	30	41
Mike	85	35
Lisa	75	87

# Selection and Projection

Name and income of persons that are less than 30:

$$\pi_{\text{name, income}}(\sigma_{\text{age} < 30}(\text{Person}))$$

```
select name, income
from   person
where  age < 30
```

name	income
Andy	21
Rob	15
Phil	30

# Naming Conventions and Renaming

To avoid ambiguities, every attribute name has two components: *RelationName.AttributeName*

When there is no ambiguity, one can drop the initial component: *RelationName*.

```
select person.name, person.income
from   person
where  person.age < 30
```

can be written as:

```
select name, income
from   person
where  age < 30
```

and also for (re-naming attributes and relations)

```
select p.name as Pname, p.income as income
from   person p
where  p.age < 30
```

## Query 1

“From the table **person**, compute a new table by selecting only the persons with an income between 20 and 30, and adding an attribute called ‘income doubled’ that has, for every tuple, double the value of **income**.”

Show the result of the query”

Person

name	age	income
------	-----	--------

## Query 1: Solution

```
select name, age, income,  
       (income * 2) as income-doubled  
from   person  
where  income >= 20 and income <= 30
```

name	age	income	income-doubled
Andy	27	21	42
Phil	26	30	<u>60</u>
Frank	60	20	<u>40</u>

# Expressions in the Target List

```
select income/4 as quarterlyIncome
from   person
where  name = 'Greg'
```

## Complex Conditions in the “where” Clause

```
select *
from   person
where  income > 25
       and (age < 30 or age > 60)
```

## The “like” Condition

The persons having a name that starts with 'A' and has a 'd' as the third letter:

```
select *
from   person
where  name like 'A_d%'
```

- ‘\_’ matches a single letter
- ‘%’ matches a string

```
[nuttdb=# select * from person where name like 'A_d%';
 name | age | income
-----+-----+-----
 Andy |  27 |      21
(1 row)
```



## Query 2

“From the table **employee**, calculate a new table by selecting only employees from the branches whose name start with ‘L’ and salary is less than 50, projecting the data on the attribute **empNo**, **salary**, **branch** and adding an attribute that has, for every tuple, twice the value of the attribute **salary**.

Show the result of the query on the following table”

Employee

empNo	surname	branch	salary
7309	Black	York	55
5998	Black	Glasgow	64
9553	Brown	London	44
5698	Brown	London	64

## Query 2

```
select empNo, branch, salary,  
       salary*2 as doubleSal  
from   employee  
where  branch like 'L%'  
       and salary < 50
```

### Employee

empNo	branch	salary	doubleSal
9553	London	44	88

## Selection, Projection, and Join

Using **select** statements with a single relation in the **from** clause we can realise:

- selections,
- projections,
- renamings

**Joins** (and Cartesian products) are realised by using two or more relations in the **from** clause

## SQL and Relational Algebra (cntd)

Given the relations:  $R1(A1,A2)$  and  $R2(A3,A4)$ ,

```
select R1.A1, R2.A4  
from   R1, R2  
where  R1.A2 = R2.A3
```

corresponds to:

$$\pi_{A1,A4} (\sigma_{A2=A3} (R1 \times R2))$$

## Query 3:

“The fathers of persons who earn more than 20K”

$\pi_{\text{father}}(\text{FatherChild} \bowtie_{\text{child=name}} \sigma_{\text{income}>20}(\text{Person}))$

```
select distinct fc.father
from   person p, fatherChild fc
where  fc.child = p.name
       and p.income > 20
```

```
[nuttdb=# select fc.father
[nuttdb=# from person p, fatherchild fc
[nuttdb=# where fc.child = p.name and p.income > 20;
 father
-----
 Greg
 Greg
 Frank
(3 rows)
```

## Query 4

“Father and mother of every person”

... can be calculated in relational algebra by means of a  
**natural join**

FatherChild ⋈ MotherChild

```
select fc.child, fc.father, mc.mother
from   motherChild mc, fatherChild fc
where  fc.child = mc.child
```

```
[nuttdb=# select fc.child, fc.father, mc.mother
[nuttdb=# from motherchild mc, fatherchild fc
[nuttdb=# where fc.child = mc.child;
 child | father | mother
-----+-----+-----
 Kim   | Greg   | Anne
 Phil  | Greg   | Anne
 Andy  | Frank  | Mary
 Rob   | Frank  | Mary
(4 rows)
```

## Query 5 Join and Other Operations

“Persons that earn more than their father,  
showing name, income, and income of the father”

Write the query in SQL

## Query 5.

“Persons that earn more than their father,  
showing name, income, and income of the father”

```
select c.name, c.income, f.income
from   person f, fatherChild fc, person c
where  f.name = fc.father and
       c.name = fc.child and
       c.income > f.income
```

name	income	income
Kim	41	40
Andy	21	20

(2 rows)



## select, with Renaming of the Result

For the persons that earn more than their father, show their name, income, and the income of the father

```
select c.name as child, c.income as income,  
       f.income as incomefather  
from   person f, fatherChild fc, person c  
where  f.name = fc.father and  
       fc.child = c.name and  
       c.income > f.income
```

child	income	incomefather
Kim	41	40
Andy	21	20
(2 rows)		

## Explicit Join

For every person, return the person, their father and their mother

```
select fatherChild.child, father, mother
from   motherChild join fatherChild on
       fatherChild.child = motherChild.child
```

```
select ...
from   Table { join Table on JoinCondition }, ...
[ where OtherCondition ]
```

# Explicit Join

For every person, return the person, their father and their mother

```
nuttdb=# select fatherChild.child, father, mother
nuttdb=# from    motherChild join fatherChild on
nuttdb=#         fatherChild.child = motherChild.child
[nuttdb=# ;
```

child	father	mother
Kim	Greg	Anne
Phil	Greg	Anne
Andy	Frank	Mary
Rob	Frank	Mary

(4 rows)

```
[nuttdb=# select fc.child, fc.father, mc.mother
[nuttdb=# from motherchild mc, fatherchild fc
[nuttdb=# where fc.child = mc.child;
```

child	father	mother
Kim	Greg	Anne
Phil	Greg	Anne
Andy	Frank	Mary
Rob	Frank	Mary

(4 rows)

## Query 5 with explicit joins

“For the persons that earn more than their father, show their name, income, and the income of the father”

```
select c.name, c.income, f.income
from   person c
       join fatherChild fc on c.name = fc.child
       join person f on fc.father = f.name
where  c.income > f.income
```

An equivalent formulation without explicit join:

```
select c.name, c.income, f.income
from   person c, fatherChild fc, person f
where  c.name = fc.child and
       fc.father = f.name and
       c.income > f.income
```

# Outer Join

“For every person, return the father and, if known, the mother”

```
select fatherChild.child, father, mother
from   fatherChild left outer join motherChild
      on fatherChild.child = motherChild.child
```

Note: “outer” is optional

```
select fatherChild.child, father, mother
from   fatherChild left join motherChild
      on fatherChild.child = motherChild.child
```

```
[nuttdb=# \e
  child | father | mother
-----+-----+-----
  Frank | Steve  |
  Kim   | Greg   | Anne
  Phil  | Greg   | Anne
  Andy  | Frank  | Mary
  Rob   | Frank  | Mary
(5 rows)
```

## Ordering the Result: order by

“Return name and income of persons under thirty, in alphabetic order of the names”

```
select name, income  
from person  
where age < 30  
order by name
```



ascending  
order

```
select name, income  
from person  
where age < 30  
order by name desc
```



descending  
order

## Ordering the Result: order by

```
select name, income
from person
where age < 30
```

name	income
Andy	21
Rob	15
Mary	42

```
select name, income
from person
where age < 30
order by name
```

name	income
Andy	21
Mary	42
Rob	15

# Aggregate Operators

Among the expressions in the target list (i.e., projection list), we can also have expressions that calculate values based on a group of tuples:

- count, minimum (min), maximum (max), average (avg), sum

**Example:** How many children has Frank?

```
select count(*) as NumFranksChildren
from   fatherChild
where  father = 'Frank'
```



## Results of count: Example

FatherChild

father	child
Steve	Frank
Greg	Kim
Greg	Phil
Frank	Andy
Frank	Rob

NumFranksChildren

2

## count and Null Values

```
select count(*)  
from person
```

Result = number of tuples  
= 4

```
select count(income)  
from person
```

Result = number of values  
different from NULL  
= 3

```
select count(distinct income)  
from person
```

Result = number of distinct  
values (excluding  
NULL)  
= 2

Person

name	age	income
Andy	27	21
Rob	25	NULL
Mary	55	21
Anne	50	35

## Aggregate Operators and Null Values

```
select avg(income) as meanIncome  
from person
```

Person

name	age	income
Andy	27	30
Rob	25	NULL
Mary	55	36
Anne	50	36

is  
ignored

meanIncome
34

# Aggregate Operators and the Projection List

An incorrect query (whose name should be returned?):

```
select name, max(income)
from person
```

The projection list has to be **homogeneous**, for example:

```
select min(age) , avg(income)
from person
```

## Aggregate Operators and Grouping

- Aggregation functions can be applied to partitions of the tuples of a relations
- To specify the partition of tuples, one uses the **group by** clause:

**group by** *attributeList*

## Aggregate Operators and Grouping

The number of children of every father.

```
select father, count(*) as NumChildren
from   fatherChild
group by father
```

### FatherChild

father	child
Steve	Frank
Greg	Kim
Greg	Phil
Frank	Andy
Frank	Rob

father	NumChildren
Steve	1
Greg	2
Frank	2

## Query 6: group by

“For each group of adult persons (age > 17) who have the same age, return the maximum income for every group and show the age”

Write the query in SQL!

Person

name	age	income
------	-----	--------

## Query 6

“For each group of adult persons who have the same age, return the maximum income for every group and show the age”

```
select age, max(income)
from    person
where   age > 17
group by age
```



# Grouping and Projection List

In a query that has a **group by** clause, **only** such attributes appear in the **group by** clause can appear in the projection list (except for aggregation functions)

**Example: Incorrect:** income of persons, grouped according to age

```
select age, income
from   person
group by age
```

The above is wrong ... because there could exist several values for the same group.

**Correct:** average income of persons, grouped by age.

```
select age, avg(income)
from   person
group by age
```

Make the attribute aggregate

## Grouping and Target List (cntd)

The syntactic restriction on the attributes in the select clause holds also for queries that would be semantically correct (i.e., for which there is only a single value of the attribute for every group).

**Example:** Fathers with their income and with the average income of their children.

**Incorrect:**

```
select fc.father, avg(c.income), f.income
from   person c join fatherChild fc on c.name=fc.child
        join person f on fc.father=f.name
group by fc.father
```

**Correct:**

```
select fc.father, avg(c.income), f.income
from   person c join fatherChild fc on c.name=fc.child
        join person f on fc.father=f.name
group by fc.father, f.income
```

# Conditions on Groups (“having” clause)

It is also possible to **filter the groups** using selection conditions.

Clearly, the selection of groups differs from the selection of the tuples in the **where** clause: the tuples form the groups.

To filter the groups, the “having clause” is used.

The having clause must appear after the “**group by**”

**Example:** “Fathers whose children have an average income greater 25.”

```
select fc.father, avg(c.income)
from   person c join fatherChild fc
      on c.name = fc.child
group by fc.father
having avg(c.income) > 25
```

## Query 7. where or having?

“Fathers whose children under age 30 have an average income greater 20”

## Query 7.

“Fathers whose children under the age of 30 have an average income greater 20”

```
select father, avg(f.income)
from   person c join fatherChild fc
      on c.name = fc.child
where  c.age < 30
group by cf.father
having avg(c.income) > 20
```

## Union, Intersection, and Difference

Within a **select** statement one cannot express unions.

An explicit construct is needed:

```
select ...  
union [all]  
select ...
```

With **union**, duplicates are eliminated  
(also those originating from projection).

With **union all** duplicates are kept.

## Positional Notation of Attributes

```
select father, child
from   fatherChild
union
select mother, child
from   motherChild
```

Which are the attribute names of the result?

Those of the first operand!

[nuttdb=# \e	
father	child
-----+-----	
Anne	Phil
Greg	Kim
Greg	Phil
Mary	Andy
Frank	Andy
Lisa	Greg
Frank	Rob
Lisa	Mary
Steve	Frank
Mary	Rob
Anne	Kim
(11 rows)	

- SQL matches attributes in the same position
- SQL renames the attributes of the second operand

## Positional Notation: Example

```
select father, child
from   fatherChild
union
select mother, child
from   motherChild
```

father	child
Anne	Phil
Greg	Kim
Greg	Phil
Mary	Andy
Frank	Andy
Lisa	Greg
Frank	Rob
Lisa	Mary
Steve	Frank
Mary	Rob
Anne	Kim

(11 rows)

```
select father, child
from   fatherChild
union
select child, mother
from   motherChild
```

father	child
Mary	Lisa
Greg	Kim
Greg	Phil
Frank	Andy
Phil	Anne
Kim	Anne
Andy	Mary
Greg	Lisa
Rob	Mary
Frank	Rob
Steve	Frank

(11 rows)



## Positional Notation (cntd)

Renaming does not change anything:

```
select father as parent, child
from   fatherChild
union
select child, mother as parent
from   motherChild
```

Correct (if we want to treat fathers and mothers as parents):

```
select father as parent, child
from   fatherChild
union
select mother as parent, child
from   motherChild
```

# Difference

```
select name
from   person
except
select child as name
from   fatherChild
```

We will see that differences can also be expressed with nested **select** statements.

# Intersection

```
select name
from   person
intersect
select child as name
from   fatherChild
```

name
Andy
Kim
Frank
Rob
Phil

(5 rows)

is equivalent to

```
select person.name
from   person, fatherChild
where  person.name = fatherChild.child
```

name
Frank
Kim
Phil
Andy
Rob

(5 rows)